

Kleene: Relational Algebra for Recursive Language-Model Calls

A planner, a budget and a replay-gated playbook for agents that write SQL

Marcus Elwin

github.com/MarcusElwin/kleene

October 2026

Technical report, version 2 (draft; v1 of 3 October 2026)

Abstract

Recursive Language Models (RLMs) let a model treat its context as a variable in a REPL and call itself on slices of it. Every published harness gives the model Python for that REPL. We argue that for this job SQL is the better language, and we present KLEENE, a Rust harness in which the model’s only action is to write CallSQL: a PostgreSQL-flavoured dialect in which model calls are scalar and table functions, tools are table functions or `CALL` statements, delegation is a lateral table call that opens a child session, search is a recursive CTE with a beam limit, and the answer is a relation. Because the work is a query, three things follow that a tool-calling loop cannot offer. A *planner* sees the cost of every call predicate and applies eight rewrite rules (cheap-first, memo dedupe, batching, proxy cascades, semi-joins for `EXISTS`, beam-limited recursion, join ordering with branch-and-bound, volatility fences) with a cost model that `EXPLAIN` shows before any money is spent. *Budgets* in calls, tokens, dollars, depth and wall clock are query semantics: a statement whose estimate exceeds what remains is refused with its plan. *Learning* is rows in DuckDB: solved SQL becomes a playbook only after a replay evaluation with the memo bypassed wins on solves and cost, and every adoption is tied to the evaluation that admitted it. We describe the algebra, the system and a benchmark harness with seven task packs run under three modes (learning, frozen, and a plain tool-calling agent on the same provider and budgets) on three models: Claude Opus 5.5, Claude Haiku 4.5 and GPT-6 Luna, 62 of 63 cells filled. Opus saturates the four original packs and the coding pack, so those separate on cost only: the playbook cuts calls per task, the replay gate costs more than the playbook saves at twenty tasks a pack, and the one-call baseline is cheapest wherever the context fits in a prompt. Luna and Haiku separate on accuracy. The playbook’s clearest wins are on the packs where the task is an aggregation or a draft: on rubric-graded memos learning beats frozen 18/20 to 8/20 on Luna and 11/20 to 8/20 on Haiku, and on long-context aggregation 20/20 to 15/20 on Luna at 1.3 calls a task against 34. Where the context fits one prompt and the question is a lookup, the plain agent is as accurate and far cheaper; where it does not (a 30 000-token log), KLEENE frozen solves 7/10 on Luna against the plain agent’s 3/10, at the price of hundreds of worker calls a task. On memo drafting the playbook hurt Opus (10 to 4), the one reversal. We report the numbers with their confounds (one run per cell, a shared memo in the first Haiku sweep, an untraced judge) and the adapter and batching fixes the runs forced.

1 Introduction

Recursive Language Models [42] change what an agent harness hands the model. Instead of a prompt that holds the whole context, the root model sees only the task; the context lives as a variable in a REPL, and the model writes code that peeks at it, partitions it, and calls `llm_query` or `rlm_query` on slices. The strategies that emerge (peek, grep, partition-and-map, summarise, solve in code) are what one would write by hand, and the paper reports large gains

on long-context aggregation at cost parity. The official harness and every reimplementations we know of give the model Python for that REPL. A reproduction study [39] found that recursion beyond depth one already degrades on current models, through format collapse and runaway sub-calls: the freedom of a general-purpose language is also the freedom to spend without a plan.

KLEENE starts from a different premise. The REPL is a relational engine and the model writes SQL. Three properties of SQL make it a better fit for exactly this job.

1. **It is declarative.** The model says *what* relation it wants. The harness decides how many calls to make, in what order, concurrently or not, and on which model. That is precisely the freedom a planner needs and that an imperative loop denies it.
2. **It has an algebra with known complexity.** Select-project-join is NP-complete in combined complexity [6], full relational algebra is PSPACE-complete [38], and `WITH RECURSIVE` reaches Datalog and beyond [8], up to Turing completeness with `UNION ALL` [9]. Every statement the model writes can be labelled with the fragment it lives in, and join ordering over call predicates is a combinatorial search the planner deliberately confronts [14].
3. **The trace is a relation.** Calls, tokens, dollars and depths land in DuckDB tables [30]. The model, the harness and the person inspect execution with the language they drive it with.

Once the work is a query, a planner, budgets and learning fall out as database problems rather than prompt engineering. This report describes the resulting system and what it has measured so far. Our contributions are:

- CallSQL, a dialect in which model calls, tools, delegation and budgets are first-class (§2), and a *call algebra* that attaches a call kind and a cost signature to any relational operator that evaluates a call (§3).
- A planner with eight rewrite rules and a cost model that mirrors the executor’s short-circuit semantics, exposed through `EXPLAIN` and used to refuse statements a budget cannot pay for (§3.1).
- A session loop with memoised calls, concurrent child sessions and a persistent store in which every learned object (playbook SQL, ratings, generator dials, refined prompt functions) is a versioned table adopted only after a replay evaluation (§4, §5).
- A benchmark harness with seven task packs, three modes that isolate the playbook and the SQL abstraction respectively, code and rubric oracles, and record-and-replay (§6); and the results of five real-model sweeps on three models, reported with their confounds (§7).

What this report does *not* claim is a benchmark win. The packs are small and mostly synthetic, every cell is one run, and Opus 5.5 saturates five of the seven packs. The numbers say where the abstraction pays for itself (aggregation over a context that does not fit, drafting against a rubric, a playbook on a stream of similar tasks) and where it does not (one-call lookups), and §9 lists what remains before a stronger claim can be made.

2 CallSQL: the language the model writes

Everything the model does in a session is one or more CallSQL statements in a single `sql` fence. The relational core is a PostgreSQL-flavoured subset: `SELECT` with CTEs and `WITH RECURSIVE`, inner and lateral joins, `EXISTS`, `IN`, scalar subqueries, grouping, ordering, set operations, `VALUES`, `CREATE TABLE ... AS`, `INSERT` and `DROP`, over the types `BOOLEAN`, `BIGINT`, `DOUBLE`, `TEXT`, `JSON` and `VECTOR`. Three-valued logic short-circuits `AND` and `OR` left to right, which the planner relies on. The core is tested differentially against DuckDB, recursive CTEs included (§7.1). On top of it sit five additions.

Context as a table. A run that is given a context starts with `ctx(ordinal, text)`, one row per paragraph. The root model sees the task, the catalog and the row count, not the text. It peeks with `SELECT ... LIMIT`, filters with `grep` or `LIKE`, partitions with `chunks`, and never has to hold the context in its prompt.

Model calls as functions. The scalar builtins `llm(prompt)` (with optional alias and effort arguments), `llm_bool(prompt)` and `llm_json(prompt, schema)` each cost one call per distinct argument tuple. `expand(text, n)` is a table builtin that returns n successors. Prompt-defined functions give a call a name, a return type, a model tier and a volatility:

```
CREATE FUNCTION mentions_hours(note TEXT) RETURNS BOOLEAN
AS PROMPT 'Does this note record hours worked? {note}'
MODEL 'proxy' BATCH 20 IMMUTABLE;

CREATE FUNCTION is_valid(c TEXT) RETURNS BOOLEAN
AS PROMPT 'Is {c} a valid answer to the task?'
MODEL 'worker' PROXY plausibility THRESHOLDS (0.2, 0.8);
```

MODEL picks a routing alias (`root`, `worker`, `proxy`, `judge`, or any alias the router knows). A prompt function is batched: the executor answers up to n distinct tuples per call, 20 unless `BATCH n` says otherwise (`BATCH 1` opts out), and the planner prices $\lceil \text{rows}/n \rceil$ calls. `PROXY` declares a cheap scorer the planner may cascade the predicate through. Bodies may also be SQL (`AS SQL (SELECT ...)`) or a shell command (`AS SHELL`); volatility defaults to `IMMUTABLE` for prompts, `STABLE` for SQL and `VOLATILE` for shell.

Tools. Pure tools are table functions usable in `FROM`: `files`, `lines`, `read`, `grep`, `chunks`, `env`, `git_log`, `git_diff`, `git_blame`. Side effects run as statements: `CALL shell(cmd)`, `CALL write_file(path, text)`, `CALL patch(path, old, new)`, `CALL web_search(q)`, and so on. A tool's output relation is rendered back to the model like any other result. Roles restrict which tools a session may use.

Delegation. `rlm(question [, context])` in a `CROSS JOIN LATERAL` opens one child session per input row at depth $d+1$ and returns (`answer`, `detail`). `CREATE AGENT` declares a role with a model, an effort, a tool list, a budget and a prompt; `spawn('role', task)` runs it. A child's catalog is restricted to its role's tools and its budget is the parent's remaining slice intersected with the role's budget.

```
CREATE TABLE parts AS
SELECT (ordinal / 20) AS part, string_agg(text, '\n') AS chunk
FROM ctx GROUP BY 1;

CREATE TABLE answers AS
SELECT p.part, r.answer
FROM parts p CROSS JOIN LATERAL rlm(
  'Total hours per project in these notes, as project: hours', p.chunk) r;
```

Finishing, budgets and plans. `FINAL FROM (query)` ends the session with that relation as the answer. `SET` adjusts effort, model, memo and the budget dimensions within a session. `EXPLAIN` prints the call plan (§3.1); `EXPLAIN ANALYZE` runs it and prints actuals beside estimates. The harness runs the planner on every statement before execution, and a statement whose estimate exceeds the remaining budget is returned to the model as a refusal carrying the plan rather than executed.

The opening example of the repository shows the shape the design is built for:

```

SELECT candidate
FROM possibilities
WHERE VERIFY(candidate)
AND NOT EXISTS (
  SELECT 1 FROM counterexamples ce
  WHERE REFUTE(candidate, ce.text));

```

VERIFY is one call per distinct candidate. The NOT EXISTS is an anti-semi-join that stops on the first refuting counterexample: the counterexample-guided loop of CEGIS and CEGAR [7, 34] written relationally. EXPLAIN says how many calls that is before they are made.

3 The call algebra

The ordinary operators are selection σ , projection π , join $\bowtie$, union $\cup$, difference, aggregation γ and the least fixpoint μ for recursive CTEs. KLEENE adds a *call kind* to any operator that evaluates a call expression (Table 1). The cost of an operator is the number of calls it implies times the per-alias token and dollar factors of the function it calls; selectivity of a boolean call predicate is a prior until the predicate has been observed, after which the sampled pass rate is used.

Operator	Meaning	Cost signature
λ_f (map-call)	scalar call f per distinct input tuple	$ \delta(R) \times \text{tokens}(f)$
κ_g (expand-call)	table call g per input tuple (LATERAL $g(\dots)$)	$ R \times \text{tokens}(g)$; output = branching factor
ρ_d (recurse)	rlm / spawn at depth d	the child's whole plan
$\sigma_{llm}, \bowtie_{llm}$	filter or join whose predicate contains a call	inherits λ_f ; $ R \times S $ calls for a join unless rewritten

Table 1: Call kinds. δ is duplicate elimination; identical prompts cost one call.

Every statement is also labelled with the fragment it lives in: CQ for select-project-join, FO for the full relational algebra, REC when a recursive term is present. The badge is a reminder that a recursive call plan is a program, not a lookup.

3.1 Rewrite rules and the cost model

The planner in `kleene-algebra` runs rules over the annotated plan until nothing changes. All eight are standard query-rewrite shapes applied to call costs.

1. **Cheap first.** $\sigma_{\text{pure}} \circ \sigma_{llm} \equiv \sigma_{llm} \circ \sigma_{\text{pure}}$: pure conjuncts run before call conjuncts, in filters and in join conditions. Because estimates mirror the executor's left-to-right short-circuit, the saving is visible in the numbers rather than assumed.
2. **Dedupe then map.** $\lambda_f(R) \equiv \lambda_f(\delta(R)) \bowtie R$, backed by a persistent memo keyed on model and prompt fingerprint and shared across sessions and runs.
3. **Batch.** λ_f over n tuples of a prompt function becomes one call with n numbered items, priced $\lceil \text{rows}/n \rceil$. Answers come back as `{"item": i, "answer": ...}` and are matched by item; the missing subset is re-asked once as a batch, and only then one call per tuple. The first runs lacked the re-ask and batched only on request, and paid for it (§7.5).
4. **Cascade.** $\sigma_{\text{oracle}}(R)$ with a declared proxy becomes $\text{score} \geq \tau_{hi} \vee (\text{score} \geq \tau_o \wedge \text{oracle})$, applied when the priced cascade is cheaper; the shape is LOTUS's [29], with thresholds declared rather than calibrated (§9).

5. **Semi-join.** [NOT] EXISTS with a call predicate becomes a semi or anti join that stops at the first (counter)example instead of materialising the cross product.
6. **Beam-limited recursion.** Inside μ , a trailing ORDER BY ... LIMIT k in the recursive term keeps the best k new rows of each round: breadth-first search with a model successor function becomes beam search of width k , and the planner simulates the fixpoint round by round to price it.
7. **Join ordering.** Dynamic programming over relation subsets with branch-and-bound, in the tradition of Selinger et al. [33] and Moerkotte and Neumann [28]: each split is priced with the predicates that first become applicable there, pure conjuncts narrowing the pairs and call conjuncts paying one call per surviving pair.
8. **Volatility fences.** Rules 1–7 move operators only across IMMUTABLE and STABLE operators. A VOLATILE operator (a shell command, a file write) is a fence: nothing crosses it, it is never deduplicated, and its input order is preserved. The planner trusts the label, so a side effect is never marked anything else.

Estimates use the store’s row counts, sampled selectivities, per-alias cost factors and pricing. EXPLAIN prints the plan as a tree with per-node rows, calls, tokens and dollars, the call kinds, the fences, the fragment badge, the rules that fired, the join-order plan space and the alternatives it priced.

Budgets as semantics. A Budget has calls, tokens, dollars, depth and wall clock. A child session receives the parent’s remaining slice intersected with its role’s budget, and its spending rolls up. Exceeding a budget cancels the statement, not the session, and returns a budget row to the model; a statement whose *estimate* already exceeds the remainder is refused before it runs. Depth defaults to a cap of two, following the reproduction study’s finding that deeper recursion degrades on current models [39]; it is a dial, not a constant.

4 System

KLEENE is a Rust workspace of eleven crates with a shared interface crate (`kleene-core`: values, schemas, the catalog, budgets, call kinds) that every other crate depends on. Figure 1 shows the path of one turn.

Execution. Operators are pull-based async streams of small columnar batches. Recursion runs by semi-naïve evaluation [3]: each round joins only the previous round’s delta against the recursive term, so calls in that term run only on new rows, and a round is one awaitable unit the UI can show as a frontier size. Model calls are concurrent under a per-session semaphore with a provider rate limiter; ordered results are preserved.

The sink. Every scalar, table, tool or child call passes through a `LiveSink` that looks the call up in the memo, and on a miss routes it to a provider alias, prices it at the router’s list rates, charges the budget, stores the response in the memo and emits a trace event. Identical prompts are therefore free the second time, across sessions and runs; this is also how a recorded run replays offline.

Sessions and delegation. Child sessions are the same loop one level deeper with a role, a budget slice and their own table namespace; a child’s `FINAL` returns to the parent as a row. Children of one lateral call run concurrently. Session state (tables, transcript, budget) lives in the store, so killing the process and running `kleene resume` continues where it stopped; checkpointing is a property of the storage, not a feature.




Figure 1: One turn of a session. The harness sends a cached system prefix (rules, catalog, budget) plus the transcript; the model replies with CallSQL; the statement is parsed, annotated, priced, checked against the budget, executed through a sink that memoises, routes and traces every call, and rendered back. The loop repeats until **FINAL**.

Model layer. Adapters speak the Anthropic Messages and OpenAI-compatible wire formats directly over HTTPS, plus an Open Responses adapter behind a feature flag; there are no vendor SDKs. A router maps aliases to models and prices, and a **ReplayProvider** serves recorded fixtures so the test suite and the benchmark harness can run without a key. The Anthropic adapter puts a cache breakpoint on the last message of every multi-turn request, not only on the system prefix, so a root session reads its transcript from cache from the third turn on; the plain baseline does the same. Both adapters rewrite model-written `llm_json` schemas into the subset their provider’s structured-output mode accepts (closed objects with every property required; a non-object root wrapped and unwrapped; typeless properties completed). Every one of these was forced by a benchmark run (§7).

Store and trace. DuckDB holds session tables, the memo, the trace (`trace_sessions`, `trace_statements`, call events), the continual loop’s tables and the benchmark rows. **kleene trace** runs plain DuckDB SQL over all of it, and the terminal UI’s `/trace` does the same from the prompt.

Processes. `kleene run`, `repl`, `explain`, `learn` and `bench` run the harness in one process. The terminal UI talks to an engine daemon over a Unix socket with a JSONL protocol with cursors, so a reconnecting client resumes from where it left off and `kleene attach` watches the same events headless. The UI is one scrolling stream: the model’s reply streams in, each fence runs as it lands, every statement shows its rows and its cost, child sessions nest under the turn that opened them.

5 Learning as tables

The continual loop (`kleene learn`) is itself a query over a `tasks` table: pick the next pending task by curriculum policy, run a session, judge **FINAL** with the task’s oracle, record outcome and cost, repeat. Two things keep it alive: generators that keep the table non-empty, and learning that changes how the next task is attempted.

Tasks and difficulty. Generators (`sat3`, `graph`, `puzzle`, `corpus`, `repo`, `statements`, `contracts`, `logbook`, `memo`) produce tasks deterministically from a hardness dial in $[0, 1]$ and a seed, each with a code oracle; the design follows Reasoning Gym [35] and the random 3-SAT phase transition [27]. Every task and every solver configuration carries a rating moved by a Bradley–Terry model [5] over solve and fail outcomes: a task that beats strong configurations is hard, a configuration that beats hard tasks is strong. The curriculum samples the pending task nearest even odds for the current solver, steps a generator’s dial up when its solve rate exceeds 0.7 and down below 0.3, in the spirit of the learnability target of Absolute Zero Reasoner [43].

What is learned. Nothing here updates weights. Each learned object is a table with a version and a provenance pointer:

- the **playbook**: SQL that solved a task, keyed by task kind, rendered into the cached prefix as examples for similar tasks and queryable by the model;
- **refined prompt functions** (`learn refine`): new versions of `VERIFY`-style functions in a ledger with before/after snapshots and rollback;
- the **cost model**: sampled selectivities of boolean call predicates that feed the planner (today per session; §9);
- **ratings** of tasks and solver configurations, and each generator’s **dial**.

The replay gate. A candidate playbook entry is adopted only after a replay evaluation: the harness re-runs a sample of earlier tasks of the same kind in two arms, with and without the candidate, with the memo bypassed, and adopts the entry only if the arm with it wins on solves and costs no more. The sample is three tasks by default, so a gate costs six task runs per solved task. Every adoption is tied in `playbook_evals` to the evaluation that admitted it, and `learn playbook` lists reverts, so a regression can be traced to the entry that caused it. The base system prompt is immutable and learned content is appended in marked sections. §7.2 shows that at twenty tasks a pack the gate spends far more than the playbook saves; it is the price of never adopting an entry on faith.

6 Benchmark methodology

The harness’s claim is specific: a model writing CallSQL over a relational engine solves the same tasks as a tool-calling agent for fewer calls and dollars, and the harness gets better over a stream of tasks because solved SQL is kept. The benchmark harness (`kleene bench`) is built to test exactly that, so every run records three things per task: was it solved, how many model calls it took, and what it cost.

6.1 Packs

A pack is a directory with a `pack.json`: tasks with text, optional context, setup commands, a workspace source, an oracle and a difficulty prior. Table 2 lists the seven shipped packs. The synthetic packs are frozen output of the loop’s generators, so two runs see the same tasks in the same order and `bench build` regenerates them byte for byte; the two long-context packs are stored lazily as (generator, dial, seed) triples and regenerated on load. The coding pack is hand-written: four small Python projects, each a three-step *episode* in which the second and third steps run in the workspace the previous step left, and a step’s setup installs the reference solution for the previous step when its hidden tests fail, so each step is measured on its own work. The external datasets these packs imitate (OOLONG [4], Terminal-Bench [26], FinanceBench [15], CUAD [13], Harvey LAB [12], SWE-bench [16]) are not redistributed; a

Harvey LAB checkout imports through `bench import-lab` with LAB’s own evaluator as the scorer of record.

Pack	<i>n</i>	What the model gets and must produce	Oracle	Stands in for
<code>terminal</code>	6	a workspace and a shell instruction; files left behind	<code>shell</code>	Terminal-Bench
<code>oolong-like</code>	20	dated meeting notes as <code>ctx</code> ; hours per project	<code>exact</code>	OOLONG aggregation
<code>finance-synthetic</code>	20	a synthetic income statement; one number	<code>number</code>	FinanceBench
<code>legal-synthetic</code>	20	a synthetic contract; the clause categories present	<code>exact</code>	CUAD / LAB
<code>logbook-hard</code>	30	a 30 000-token work log with corrections and distractors; one of five aggregate questions	<code>exact</code>	OOLONG at full size
<code>memo-rubric</code>	20	a 40-section agreement with an amendment and a rejected proposal; a short memo	<code>judge</code>	LAB drafting
<code>coding</code>	12	a Python project with visible tests; implement, extend, fix a bug	<code>shell</code>	SWE-bench episodes

Table 2: The seven shipped packs. `exact` compares row multisets, `number` the first cell within a tolerance, `shell` runs a command (here the grader’s hidden tests) and `judge` asks a separate model for PASS against a rubric that names every required fact and forbids the superseded values.

6.2 Modes

Every pack runs under one of three modes, and two comparisons fall out of them.

learning KLEENE: the model writes CallSQL, the playbook is shown, candidates are recorded and adopted through the gate, ratings move.

frozen KLEENE with the same model, tools and budgets, but no playbook shown, no candidate recorded, no rating moved. The control for learning.

plain a tool-calling agent on the same provider, tools and budgets, in the ReAct shape [41]: the tools go to the model as native tool definitions, every tool-use block of a reply runs and its result comes back as a tool result, and a `final` tool carries the answer. The baseline for cost parity, as in the RLM paper’s framing: same model, same spend, only the abstraction differs.

learning against **frozen** isolates the playbook; **frozen** against **plain** isolates the SQL abstraction. Two caveats matter for reading the results. Both KLEENE modes read the same store, so a frozen run on a store that has seen a learning run (or the reverse) is served from the memo for every prompt already sent; a clean comparison needs a fresh store per pack and mode, which every sweep after the first Haiku one does (§7.3). And the first Haiku sweep’s plain agent parsed one JSON action per reply from text; it was replaced by the native tool-calling agent above before the later sweeps, and the coding pack’s tasks now carry a `check` that refuses a FINAL while the project’s tests fail.

6.3 What is recorded

`bench run` writes one row per task to `evals`: run, pack, mode, position in the stream, task, solved, the oracle’s reason, calls, tokens, dollars, deepest child, turns and wall clock. `bench`

report aggregates by pack and mode; **bench csv** exports the rows; **bench curve** prints the rolling solve rate of one run; **bench plot** draws five standalone SVGs (learning curve, accuracy at cost parity, calls against difficulty, planner estimate against actual, plan-space size against query shape) without a plotting library. The replay evaluations a learning run pays for are recorded in **playbook_evals** with their own dollars and are *not* in **evals**; the true cost of a learning run is the sum. **bench run -record** saves every model reply as a fixture and **-replay** serves it back, so a run reproduces offline; **-resume** continues a run that a provider error stopped.

7 Results

7.1 Deterministic measurements

Table 3 is the part of the evidence that needs no model: the test suite runs it in CI with scripted providers, code oracles and replay fixtures.

Property	Test	Result
Relational core equals DuckDB	property-based differential suite	passes, recursive CTEs included
Join ordering, two call predicates over three relations	estimate and actual calls against a counting sink	as written 12 000 estimated calls, chosen order 60; actual $\geq 10\times$ fewer with identical rows
Plan-space growth	seven-way join	665 280 bushy orders priced through 1 623 subset splits (309 pruned)
Cascade	proxy at $0.1\times$ cost, thresholds (0.2, 0.8)	100 proxy + 60 oracle calls beat 100 oracle calls by estimate; on data the oracle is asked only for the band (5 of 20)
Beam recursion	expand $\times 4$, three rounds	20 calls with beam 5 against > 100 unbounded
Continual loop	scripted solver over generated puzzles	tasks solved, playbook v1 adopted after a replay eval, ratings and dial move, restart resumes
Terminal-Bench-style pack	shell tasks from SQL	tasks solved, oracles run in fresh workspaces
Plain-agent baseline	same provider, native tool calls	tool call then final; the control for cost parity

Table 3: What the deterministic suite establishes.

7.2 Claude Opus 5.5 on the first four packs

The four original packs ran once each under the three modes on 27 September 2026 with the default routing: **root** on **claude-opus-5-5** at high effort, **worker** and **judge** on **claude-sonnet-5**, **proxy** on **claude-haiku-4-5**, priced at the router’s list rates. Table 4 gives the per-pack rows; Figure 2 the plots **bench plot** drew from the same store. **dollars** is the pack total, calls and tokens are per task, and the learning column excludes the gate.

Accuracy is saturated. Opus 5.5 solves every task of every pack in every mode. The packs were sized for a weaker model at dial 0.5; they separate nothing on accuracy, so what follows is about calls and dollars.

The playbook cuts calls per task. Learning against frozen: terminal 2.17 to 1.33, oolong 4.90 to 3.85, finance 3.45 to 2.40, legal 12.65 to 10.60. Only two packs adopted anything, though:

Pack	Mode	Solved	Calls / task	Tokens / task	Dollars
terminal (6)	frozen	6/6	2.17	10 551	0.07
	learning	6/6	1.33	6 742	0.04
	plain	6/6	2.50	2 395	0.03
oolong-like (20)	frozen	20/20	4.90	30 778	1.51
	learning	20/20	3.85	41 991	1.92
	plain	20/20	1.00	3 084	0.39
finance-synthetic (20)	frozen	20/20	3.45	42 177	0.83
	learning	20/20	2.40	29 999	0.63
	plain	20/20	1.00	1 244	0.07
legal-synthetic (20)	frozen	20/20	12.65	62 452	1.69
	learning	20/20	10.60	58 315	2.61
	plain	20/20	1.00	1 475	0.11

Table 4: Claude Opus 5.5, 27 September 2026, one run per cell. Gate replays (not included): \$1.25 terminal, \$10.70 oolong, \$5.46 finance, \$16.63 legal.


oolong (16 of 20 candidates) and legal (5 of 21). On terminal and finance every candidate was rejected because its replay cost more than the baseline’s, so those learning runs are frozen runs with ratings on, and their call reduction is run-to-run variance, not learning. On oolong the second half of the learning run took 3.9 calls per task where the first half took 3.8: most of the gain arrived with the first adopted entries. On legal the halves are 13.7 and 7.5, the curve the design predicts.

The gate is the cost. Each candidate is replayed on three earlier tasks in both arms, six task runs per solved task, and those replays cost \$1.25, \$10.70, \$5.46 and \$16.63 on the four packs against \$0.04, \$1.92, \$0.63 and \$2.61 for the counted runs. At twenty tasks a pack the gate spends far more than the playbook saves. It earns its keep only over a long stream, and a cheaper gate (a smaller sample, a proxy model, a cached baseline arm) is the obvious next change.


The plain agent wins on cost here. Every context in these packs fits in one prompt (under 4000 characters), so the baseline answers in one call with no tool use: \$0.39 against \$1.51 on oolong, \$0.07 against \$0.83 on finance, \$0.11 against \$1.69 on legal. The SQL abstraction pays for a catalog, a planner and a transcript of rendered results on every turn, and at this scale that is pure overhead. The claim of §1 is about work that does not fit in one call; these packs do not test it, and the numbers say so.

Per-clause calls are cheap and visible. On legal the model classified clause by clause with a prompt function on the `worker` alias, up to 26 calls a task at about \$0.10: the λ_f map-call doing what it is for, many small calls on a cheap model, priced before they run.

Costs drift up within a run. Per-task dollars on oolong learning rose from \$0.084 in the first half to \$0.107 in the second while calls stayed flat. Two things grow the prompt: the adopted playbook entries, and, as it turned out, the scratch tables every session created in the shared store, which the catalog listed to every later session (287 such tables after these runs). The second was a bug: it also changed the prompt fingerprint from task to task, so a recorded run could not be replayed from its own fixtures. Bench sessions now drop their tables when they finish.



(a) Cumulative accuracy against cumulative dollars.



(b) Calls per task against the task’s rating.

Figure 2: The Opus 5.5 run. Every line reaches accuracy 1; the plot separates the modes on cost only, with the plain agent leftmost on every pack.


7.3 Claude Haiku 4.5 on the harder packs

The harder packs ran on 1 October 2026 with every solver alias on `claude-haiku-4-5-20251001` and the `judge` alias on `claude-sonnet-5-5`. The sweep was stopped before it finished; Table 5 is what completed. Not in the table: `logbook-hard` in any mode (a frozen task took about 500 calls and over twenty minutes at Haiku’s pace, so the pack was capped at ten tasks and then stopped before the first row landed), and `coding` in plain mode (below). Gate replays cost \$1.43 on memo-rubric (11 candidates, 3 adopted) and \$1.04 on coding (2 candidates, 1 adopted); the counted runs cost \$3.36 in all.


These packs separate on accuracy. Where Opus 5.5 solved every task of the first four packs, Haiku solves 17% of the coding steps and 40 to 55% of the memos. The coding steps it solves are the first step of an episode; it solved no “extend” or “fix the bug report” step in either KLEENE mode.

Pack	Mode	Solved	Calls / task	Tokens / task	Dollars
coding (12)	frozen	2/12	4.58	48 375	0.87
	learning	2/12	2.67	29 910	0.47
memo-rubric (20)	frozen	8/20	19.95	49 043	1.30
	learning	11/20	3.80	23 041	0.37
	plain	8/20	3.15	10 764	0.35

Table 5: Claude Haiku 4.5 solvers with a Sonnet 5.5 judge, 1 October 2026, one run per cell; the sweep was stopped before `logbook-hard` and the coding plain run produced rows. The learning rows share the frozen rows’ memo (see text).



(a) Rolling-5 accuracy by position in the stream.



(b) Cumulative accuracy against cumulative dollars.

Figure 3: The Haiku 4.5 run on `memo-rubric` and `coding`. On `memo-rubric` the learning line ends above frozen and plain at a fraction of frozen’s spend; on `coding` both KLEENE lines fall to zero after the first step of each episode.

Learning beat frozen on memo-rubric, on both axes. 11/20 against 8/20, at 3.8 calls a task against 19.95 and \$0.37 against \$1.30. Two effects are mixed in. Three playbook entries were adopted, and the second half of the learning run solved 5 of 10 where the frozen run’s first half solved 3 of 10. But most of the frozen run’s calls are a failure mode the playbook happens to steer around: Haiku writes `llm_json` output schemas the API rejects (`additionalProperties` unset, `minItems` above 1, a bare `{"hours": "number"}`), and the statement retries until the call budget is gone; five frozen tasks took 22 to 141 calls each. The adapter now rewrites a model-written schema into the accepted subset before the request; these runs predate it.

Learning and frozen share the memo. The learning run came second in the same store, so a prompt identical to one the frozen run sent was answered from the memo for free. Two coding steps in the learning run finished with zero calls: they replayed the frozen run’s failed answers. The learning column’s calls and dollars are therefore not an independent measurement; every later sweep used a fresh store per pack and mode.

The plain baseline needs a protocol the model keeps. The agent asks for one JSON action per reply. Haiku answers with five to ten actions and invented tool output between them; the first extractor took the span from the first brace to the last, which never parsed, so on every coding task the model saw only a parse error and gave up after two turns (0/12 at \$0.45). The extractor now takes the first balanced object and runs it. With that, Haiku worked through the first coding task for all 30 turns without solving it (\$1.11) and the rerun was stopped before a row landed. On memo-rubric, where the task fits in one call, the plain agent matches frozen on accuracy at a quarter of the cost, and learning beats it by three tasks for the same money.

The judge wobbles. Sonnet 5.5’s first lines included “PASS (with caveat)”, “PASS (provisional)” and “PASS or FAIL?”, all counted as passes by the first parser. It now requires a first line that says PASS and does not mention FAIL. The memo frozen and learning runs predate that change; the plain run does not.

Two provider facts. Haiku rejects the `effort` parameter the adapter sent along with the rest of the model family, so under the default routing every `proxy` call failed until the adapter learned to omit it; and `claude-sonnet-5-5` was served on 1 October where it was not on 27 September.

7.4 GPT-6 Luna on all seven packs

All seven packs ran on 5 October 2026 with every solver alias on `gpt-6-luna` (OpenAI chat completions, default reasoning effort) and the `judge` alias on `claude-sonnet-5-5`. Luna is priced at \$0.10 per million input tokens and \$0.50 per million output, a tenth of Haiku and a fortieth of Opus. Each pack and mode ran in its own store, so the learning columns are independent of the frozen ones. `logbook-hard` was capped at ten tasks after the Haiku experience; its learning run was stopped after three tasks, when the second task alone took 1 705 calls and ran out of turns. Table 6 gives the rows; Figure 4 the learning curves. The counted runs cost \$0.74 of Luna in all; the gate’s replays (84 candidates, 17 adopted) another \$0.31; the Sonnet judge’s calls are not traced, roughly a dollar at most.

The original packs separate on accuracy with Luna. Opus solved every task of these packs; Luna does not, so they measure something now. The one-call plain agent is perfect on finance and legal, where the context fits and the question is one lookup, and weakest on oolong-like (14/20), where aggregation over the whole context is the task. Learning on oolong-like

Pack	Mode	Solved	Calls / task	Tokens / task	Dollars
terminal (6)	frozen	6/6	1.50	4 128	0.00
	learning	5/6	1.50	4 219	0.00
	plain	6/6	2.33	1 584	0.00
oolong-like (20)	frozen	15/20	33.95	19 230	0.04
	learning	20/20	1.30	4 516	0.00
	plain	14/20	1.00	2 031	0.01
finance-synthetic (20)	frozen	14/20	4.15	11 789	0.01
	learning	15/20	3.45	12 081	0.01
	plain	20/20	1.00	950	0.00
legal-synthetic (20)	frozen	18/20	15.65	23 973	0.03
	learning	15/20	4.25	15 893	0.01
	plain	20/20	1.00	1 004	0.00
coding (12)	frozen	11/12	12.08	55 286	0.02
	learning	10/12	8.17	42 094	0.02
	plain	9/12	10.25	27 860	0.03
memo-rubric (20)	frozen	8/20	13.60	36 571	0.03
	learning	18/20	6.10	27 683	0.02
	plain	11/20	1.00	1 978	0.01
logbook-hard (10)	frozen	7/10	247.30	180 736	0.24
logbook-hard (3 of 10)	learning	1/3	686.67	746 259	0.19
logbook-hard (10)	plain	3/10	2.30	97 000	0.07

Table 6: GPT-6 Luna solvers with a Sonnet 5.5 judge, 5 October 2026, one run per cell, one store per cell. Dollars below half a cent print as 0.00.




Figure 4: Rolling-5 accuracy by position in the stream, Luna run, all packs and modes. The memo-rubric learning line sits at or above 0.8 from the start where frozen hovers near 0.4; the coding lines stay high in every mode.

solved 20/20 at 1.3 calls a task against frozen’s 15/20 at 34: the adopted playbook entry carries the aggregation query, and the rest is a lookup.

Luna solves what Haiku could not, for a tenth of the money. Eleven of the twelve coding steps frozen against Haiku’s two, including every “extend” and “fix the bug report” step but one; seven of ten logbook tasks frozen where Haiku had not finished one. Each coding step cost about a fifth of a cent.

Learning beat frozen on memo-rubric again, now cleanly. 18/20 against 8/20 at 6.1 calls a task against 13.6, in separate stores, so no memo replay hides in the learning column. Four memo entries were adopted out of 18 candidates; both halves of the learning run solved 9 of 10, so the gap to frozen is there from the start and is mostly the shorter route the playbook gives. On coding, learning cost fewer calls (8.2 against 12.1) and lost one task.

Plain is the cheap mode where the context fits, and no further. On memo-rubric the one-call agent solves 11/20 for half a cent; on logbook-hard, whose contexts are around 30 000 tokens, it gets 3/10, and on coding 9/12 with a tool loop. Frozen beats it on every one of the harder packs on accuracy.

Logbook-hard is where the calls go. Frozen tasks took 113 to 719 calls each, nearly all of them `worker` calls from batched `llm_json` extraction over the log: Luna sometimes answers a batch of 20 with the wrong number of elements, and the harness then retried one call per row. Learning was worse, not better: the one task it finished past the first took 1 705 calls and no `FINAL`. This per-row fallback was the next thing fixed (§3.1, rule 3).

Three provider facts. OpenAI rejects `max_tokens` for GPT-5 and GPT-6 models (the adapter sends `max_completion_tokens`); strict structured outputs want an object root with every property required and `additionalProperties: false`, so a non-object schema is wrapped and unwrapped; and a model-written schema with a typeless property or a partial `required` list is rejected outright, so the adapter completes it. Before those fixes a Luna run could not make a single `llm_json` call.

7.5 Claude Haiku 4.5, the remaining cells

The cells the 1 October sweep left empty ran on 6 October 2026 with the same routing, one fresh store per pack and mode, after two adapter changes: the conversation is cached turn by turn, and the structured-output schema drops `maxItems`, which the API rejects and every batched `llm_json` call used to set (Table 7). The counted runs cost \$20.07, of which `logbook-hard` took \$15.66; the gate’s replays (62 candidates, 22 adopted) another \$5.64.

On the original packs Haiku is close to Luna, at ten times the price. Learning beats frozen on oolong-like (19 against 14 at 2 calls a task against 50) and legal (17 against 14 at 10 against 44), as it did for Luna; the one-call plain agent is perfect on finance and oolong-like.

Logbook-hard is a worker-call storm. The frozen run made 79 324 `worker` calls of about 37 tokens each, \$4.30 of its \$5.13: Haiku writes a scalar `llm()` per row over the 30 000-token log, and when it does batch it answers with the wrong number of elements two times in three, so the harness fell back to one call per row. Two tasks took 5 000 and 8 500 calls. The root session itself was cheap (121 calls, \$0.83, 1.2M tokens read from cache against 36 000 billed in full). This is the fallback the Luna run flagged, now with a price on it; the fix landed the same

Pack	Mode	Solved	Calls / task	Tokens / task	Dollars
terminal (6)	frozen	6/6	3.00	11 190	0.07
	learning	6/6	2.17	8 159	0.05
	plain	6/6	2.50	7 286	0.05
oolong-like (20)	frozen	14/20	50.40	53 217	0.75
	learning	19/20	1.95	9 115	0.20
	plain	20/20	4.15	24 932	0.49
finance-synthetic (20)	frozen	20/20	7.25	30 274	0.43
	learning	19/20	7.60	18 963	0.29
	plain	20/20	1.65	5 375	0.13
legal-synthetic (20)	frozen	14/20	43.50	48 594	0.72
	learning	17/20	9.95	16 541	0.29
	plain	15/20	1.85	6 227	0.16
coding (12)	plain	4/12	28.58	218 019	0.78
logbook-hard (10)	frozen	1/10	1 702.10	476 213	5.13
	learning	1/10	1 694.40	591 443	5.95
	plain	2/10	22.20	1 365 343	4.58

Table 7: Claude Haiku 4.5 solvers with a Sonnet 5.5 judge, 6 October 2026, one run per cell, one store per cell. The coding plain row uses the native tool-calling agent; the coding frozen and learning rows and the memo-rubric rows are those of Table 5.

day: prompt functions batch 20 rows per call unless told otherwise, and a short batch answer is re-asked by item instead of one call per row.

Caching works where the prompt is long enough. Haiku’s minimum cacheable prompt is 4096 tokens, so terminal and the per-row worker calls never cache; everywhere else the root session reads its history from cache from the third turn on.

7.6 Claude Opus 5.5 on coding and memo-rubric

The two hard packs Opus had not run yet, in all three modes, with the routing of the September run (**root** on Opus 5.5 at high effort, **worker** and **judge** on Sonnet 5.5, **proxy** on Haiku 4.5), one fresh store per pack and mode, the conversation cached and prompt functions batched by default, on 6 October 2026 (Table 8). The counted runs cost \$7.38; the gate’s replays (16 candidates, 8 adopted) another \$9.22, \$7.67 of it on coding alone, so a learning run of coding costs five times what its rows show. **logbook-hard** has not run on Opus; it is the one empty cell of Table 9.

Pack	Mode	Solved	Calls / task	Tokens / task	Dollars
coding (12)	frozen	12/12	8.58	76 533	1.31
	learning	12/12	8.00	88 984	1.70
	plain	12/12	9.42	68 156	1.32
memo-rubric (20)	frozen	10/20	5.30	40 086	1.26
	learning	4/20	5.35	44 064	1.24
	plain	10/20	1.00	5 763	0.55

Table 8: Claude Opus 5.5 root with Sonnet 5.5 workers and judge, 6 October 2026, one run per cell, one store per cell.

Coding is saturated for Opus, as the original packs were. Every task in every mode, where Luna solved 11, 10 and 9 and Haiku 2, 2 and 4. The three modes cost the same within a few cents, so the pack no longer separates modes at this model size.

Memo-rubric is the one pack where Opus is behind Luna. Half the tasks in frozen and plain against Luna’s 8 and 11, and learning made it worse: the adopted entries pulled Opus from 10 to 4 solved at the same cost, where Luna’s learning run on the same pack went from 8 to 18. The judge is Sonnet 5.5 in both runs, so the difference is in what the models write, not in how it is scored. The gate admits an entry that wins its replay sample of three; an entry that wins three memos and loses six is the failure mode the sample size allows, and this is the first run to show it.

The caching and batching changes hold at Opus prices. Learning on coding read 5.0M tokens from cache against 30 000 billed in full, and the batched prompt functions kept the worker side to a few hundred calls per run where the Haiku logbook-hard runs made tens of thousands.

7.7 Across models

Table 9 collects the pass rates of every cell, and Figure 5 plots pass rate against cost per task for the three packs that separate on accuracy. Reading across:

Pack	n	Claude Opus 5.5	Claude Haiku 4.5	GPT-6 Luna
terminal	6	100 / 100 / 100	100 / 100 / 100	83 / 100 / 100
oolong-like	20	100 / 100 / 100	95 / 70 / 100	100 / 75 / 70
finance-synthetic	20	100 / 100 / 100	95 / 100 / 100	75 / 70 / 100
legal-synthetic	20	100 / 100 / 100	85 / 70 / 75	75 / 90 / 100
coding	12	100 / 100 / 100	17 / 17 / 33	83 / 92 / 75
memo-rubric	20	20 / 50 / 50	55 / 40 / 40	90 / 40 / 55
logbook-hard	10	pending	10 / 10 / 20	33 / 70 / 30

Table 9: Pass rate in percent, one column per model; each cell reads learning / frozen / plain with the best mode in bold. The Haiku coding frozen and learning cells and all Haiku memo-rubric cells are from the 1 October sweep (shared memo, text-protocol baseline); every other cell ran in its own store. The Luna logbook-hard learning cell is 1 of 3 tasks. Opus on logbook-hard is pending.

- **Where the task is an aggregation or a draft, learning wins.** On oolong-like and memo-rubric the learning mode is the most accurate mode for Luna and for Haiku, and on oolong-like it is also the cheapest KLEENE mode by an order of magnitude in calls. The adopted entries carry the shape of the answer (the grouping query, the memo’s structure), so the stream becomes lookups.
- **Where the task is a lookup, plain wins.** On finance and legal the one-call agent is perfect for Luna and Opus and within a task of it for Haiku, at a tenth to a hundredth of the cost. The SQL abstraction pays for a catalog, a planner and rendered results on every turn, and a lookup does not need them.
- **Where the context does not fit, frozen wins and pays for it.** logbook-hard is the pack built for the claim of §1: Luna frozen solves 7/10 where the plain agent solves 3/10, and Haiku frozen solves 1/10 at 1 700 calls a task. The abstraction gets the answer when the model can write the extraction; the cost is the worker-call volume, which the batching change has since cut and which the pending Opus run will measure.




Figure 5: Pass rate against mean dollars per task of the solver’s own calls, every model and mode, for the three packs that separate on accuracy. The dashed line is the Pareto frontier. Gate replays and judge calls are not in the cost.

- **Stronger models erase the modes.** Opus solves every coding step in every mode at the same price, as it solved the original packs. A pack that separates modes at one model size is saturated at the next, so the benchmark set has to grow with the models.
- **The gate is still the cost of learning.** Across the three sweeps the replays cost between 40% and 125% of the counted runs, and on Opus coding five times.

7.8 Planner estimates against actuals

The Haiku run was the first made after `bench run` began tracing statements, so its store holds an estimate and an actual call count for each of the 1 028 statements the planner priced (Figure 6). Reading the plot: 670 statements made no calls and were estimated at none; 254 were estimated at one call and made one; 82 were estimated at one and made none, which is the memo answering a repeated prompt. The planner never under-estimated by more than a call. Where it was wrong it was high, by up to two orders of magnitude on about a dozen statements, each a per-row call priced on a table’s row count that then failed early or was served from the memo. For budget refusal that is the safe direction; for the learned cost model it is the error to drive down. The later sweeps traced 4 517 (Luna), 2 755 (Haiku, 6 October) and 928 (Opus, 6 October) further statements, with the same shape.




Figure 6: Estimated against actual calls per statement, Haiku run. Points on the diagonal are exact; the outliers lie below it (over-estimates).

7.9 What the numbers say

Taken together, the sweeps support four statements and refuse a fifth. First, the planner, the budget and the trace work as designed on real models of three families: estimates are conservative, refusals carry plans, per-row calls land on the cheap alias, and every cost in this section was read from the store. Second, the playbook reduces calls per task when something is adopted, and on the packs whose tasks share a shape (aggregation, drafting against a rubric) it also raises solves, for two of the three models. Third, the replay gate, at its current sample size, is the dominant cost of learning at pack scale, and a sample of three can admit an entry that hurts (Opus on memo-rubric). Fourth, the long-context case the design was built for is now measured: on the one pack whose context does not fit a prompt, KLEENE frozen is the most accurate mode for both models that ran it, at a call volume the batching change has since reduced. What the runs still do not support is a general cost-parity claim: wherever the context fits in one prompt and the question is a lookup, the plain agent is as accurate and far cheaper, and the strongest model saturates five of the seven packs.

8 Related work

Recursive language models. Zhang et al. [42] introduced the REPL-variable framing and the `llm_query` / `rlm_query` sub-calls that KLEENE maps onto `llm` and `rlm`; their evaluation suite (OOLONG, BrowseComp-Plus) is the headline comparison the **logbook-hard** pack is built toward. Wang [39] reproduced the result and documented the depth-two degradation that set KLEENE’s default depth cap. DSPy’s `dspy.RLM` [18] is the most used implementation; like the reference harness it is Python-based with fixed iteration and call caps rather than a priced plan. CodeAct [40] made the general case that executable code is a better action space than JSON tool calls; KLEENE agrees and narrows the language to one with an optimizer.

SQL and relational operators over language models. The closest systems embed model calls in a relational language and optimise them. LOTUS [29] defines semantic filter, map, join, aggregate and top-*k* operators with proxy cascades and statistical guarantees; KLEENE’s cascade rule has LOTUS’s shape with thresholds declared rather than calibrated. Palimpsest [23] and Abacus [31] run a Cascades-style [11] cost-based optimiser over model and prompt choices using sampled sentinel plans; KLEENE’s planner is simpler (eight rules, one cost model) but operates on the model’s own SQL at run time rather than on a pre-declared pipeline. ThalamusDB [17] adds `NLfilter` and `NLjoin` to DuckDB with progressive error bounds and call budgets. BlendSQL [10]

is closest in spirit to CallSQL: a dialect with LLMMap, LLMQA and LLMJoin whose AST is rewritten so cheap SQL predicates run first; SUQL [24] exposes free-text primitives as lazily evaluated Postgres UDFs; Galois [32] treats the model itself as the database executing plan fragments; CAESURA [37] uses the model as the planner; Evaporate [2] synthesises extraction code instead of prompting per row, which is what a refined `CREATE FUNCTION ... AS SQL` amounts to. Liu et al. [25] reorder rows to share prompt prefixes in relational workloads, a rule KLEENE does not yet have. Scallop [22] makes foundation models foreign predicates in Datalog. Su et al. [36] give a taxonomy and benchmark of model-enhanced relational operators. What distinguishes KLEENE from all of these is that the *model writes the query*, turn by turn, with delegation, tools, budgets and learning in the same language, rather than a person writing a pipeline the model fills in.

Query optimisation. The planner’s machinery is classical: System R’s dynamic programming [33], bushy enumeration without cross products [28], semi-naive fixpoint evaluation [3], and the complexity results of Chandra and Merlin [6], Vardi [38], Aho and Ullman [1] and Dantsin et al. [8] that label each statement’s fragment. The name honours three of Kleene’s results [19–21]: the star as closure under repetition (what a recursive CTE computes), the fixed-point construction (how it computes it), and the recursion theorem (a session opening a child session is the same loop referring to itself).

Continual harnesses and curricula. The loop’s learnability target follows Absolute Zero Reasoner [43]; its generators follow Reasoning Gym [35] and the 3-SAT phase transition [27]; its ratings are Bradley–Terry [5]. The replay gate is the harness-level analogue of the default-FAIL test contract common to long-running agent harnesses: nothing is adopted that has not beaten the current version on held tasks.

9 Limitations and future work

- **One run per cell.** Every number in §7 is a single run. On twenty tasks accuracy moves in steps of five points, and a difference of one task is not a result. Repeated runs are the first thing to add.
- **Saturated and synthetic packs.** Opus 5.5 saturates five of the seven packs, and the synthetic packs are generated, so a solver that learns the generator’s shape will look better than one meeting the real datasets. Importers exist for the real OOLONG (from Hugging Face), Harvey LAB and two contract-redlining datasets, but no model has run them yet.
- **One cell is pending.** `logbook-hard` has not run on Opus 5.5; with the batching default in place it is the run that will say what the long-context case costs on the strongest model.
- **The first Haiku sweep is not like the others.** Its coding and memo-rubric cells share a memo between frozen and learning, use the text protocol baseline, and predate the finish check and the schema rewrite. They are kept because they are what was measured; a rerun would replace them.
- **The gate’s cost and its sample.** Six task runs per solved task dominates a learning run at pack scale, and a sample of three admitted entries that hurt Opus on memo-rubric. A larger or adaptive sample, a proxy model for the replay, and a cached baseline arm are the candidates.
- **Judge calls are untraced.** The Sonnet judge’s cost is not in any table; it is one short call per memo task, roughly a dollar a sweep.

- **Declared thresholds, per-session cost model.** Cascade thresholds are declared, not calibrated from a sample as in LOTUS; sampled selectivities persist only within a session rather than as versioned tables.
- **Schema rewriting is partial.** Model-written `llm_json` schemas are rewritten into each provider’s accepted subset only for the shapes seen so far.

Beyond these, the design names measurements the harness supports but no run has yet exercised: transfer (learn on a prefix with `-limit`, run the rest frozen on the same store, compare against a never-learned store), reverts through `learn playbook`, and the plan-space plot over the shapes models actually write.

10 Reproducibility

KLEENE v0.1.0 (6 October 2026) ships prebuilt binaries for macOS and Linux; the repository is its own Homebrew tap:

```
curl -fsSL https://raw.githubusercontent.com/MarcusElwin/kleene/main/install.sh | sh
brew tap MarcusElwin/kleene https://github.com/MarcusElwin/kleene && brew install
kleene
```

The repository is a Rust workspace; from source, DuckDB builds once (about ten minutes). Every claim in Table 3 is a test:

```
cargo test --workspace --all-features
cargo run -- repl < demos/planner/three_way.sql      # the plan space, live
cargo run -- learn run --tasks 20 --generators puzzle,corpus
cargo run -- bench run tasks/terminal --mode frozen --record fixtures/terminal
cargo run -- bench report
cargo run -- bench plot plots/
```

The per-task rows of every sweep are under `plots/evals.csv` for §7.2, and the directories `haiku-2026-10-01`, `luna-2026-10-05`, `haiku-2026-10-06` and `opus-2026-10-06` for the others, each with the `bench report` output, the router file and the SVGs beside it. `kleene bench results` regenerates Table 9 and Figure 5 from those CSVs. The model replies of the first Opus run were recorded but are not checked in, because the catalog bug of §7.2 means they cannot replay the run; later runs record with `-record` and replay with `-replay`. Prices are the router’s list rates at the time of each run, per million input / output tokens: \$4 / \$20 for Opus 5.5, \$2 / \$10 for Sonnet 5.5, \$1 / \$5 for Haiku 4.5, \$0.10 / \$0.50 for GPT-6 Luna.

Acknowledgements. The implementation, the benchmark runs and much of the documentation were produced with coding agents working from the plan in `docs/PLAN.md`; the design decisions, the measurements and the mistakes are the author’s.

References

- [1] Alfred V. Aho and Jeffrey D. Ullman. Universality of data retrieval languages. In *Proceedings of the 6th ACM SIGACT-SIGPLAN Symposium on Principles of Programming Languages (POPL)*, pages 110–119, 1979. doi: 10.1145/567752.567763.
- [2] Simran Arora, Brandon Yang, Sabri Eyuboglu, Avaniika Narayan, Andrew Hojel, Immanuel Trummer, and Christopher Ré. Language models enable simple systems for generating structured views of heterogeneous data lakes. *arXiv preprint arXiv:2304.09433*, 2023. URL <https://arxiv.org/abs/2304.09433>.

- [3] François Bancilhon and Raghu Ramakrishnan. An amateur’s introduction to recursive query processing strategies. In *Proceedings of the 1986 ACM SIGMOD International Conference on Management of Data*, pages 16–52, 1986. doi: 10.1145/16894.16859.
- [4] Amanda Bertsch, Adithya Pratapa, Teruko Mitamura, and Graham Neubig. Oolong: Evaluating long context reasoning and aggregation capabilities. *arXiv preprint arXiv:2511.02817*, 2025. URL <https://arxiv.org/abs/2511.02817>.
- [5] Ralph Allan Bradley and Milton E. Terry. Rank analysis of incomplete block designs: I. the method of paired comparisons. *Biometrika*, 39(3/4):324–345, 1952. doi: 10.2307/2334029.
- [6] Ashok K. Chandra and Philip M. Merlin. Optimal implementation of conjunctive queries in relational data bases. In *Proceedings of the 9th Annual ACM Symposium on Theory of Computing (STOC)*, pages 77–90, 1977. doi: 10.1145/800105.803397.
- [7] Edmund Clarke, Orna Grumberg, Somesh Jha, Yuan Lu, and Helmut Veith. Counterexample-guided abstraction refinement. In *Computer Aided Verification (CAV)*, volume 1855 of *Lecture Notes in Computer Science*, pages 154–169, 2000. doi: 10.1007/10722167_15.
- [8] Evgeny Dantsin, Thomas Eiter, Georg Gottlob, and Andrei Voronkov. Complexity and expressive power of logic programming. *ACM Computing Surveys*, 33(3):374–425, 2001. doi: 10.1145/502807.502810.
- [9] Andrew Gierth. Cyclic tag system. PostgreSQL wiki, https://wiki.postgresql.org/wiki/Cyclic_Tag_System, 2010. A Turing-complete construction in SQL:2003 with WITH RECURSIVE.
- [10] Parker Glenn, Parag Pravin Dakle, Liang Wang, and Preethi Raghavan. BlendSQL: A scalable dialect for unifying hybrid question answering in relational algebra. *arXiv preprint arXiv:2402.17882*, 2024. URL <https://arxiv.org/abs/2402.17882>.
- [11] Goetz Graefe. The cascades framework for query optimization. *IEEE Data Engineering Bulletin*, 18(3):19–29, 1995.
- [12] Harvey AI. Harvey LAB: the legal agent benchmark. <https://github.com/harveyai/harvey-labs>, 2026. Repository and evaluator; accessed 2026-10-03.
- [13] Dan Hendrycks, Collin Burns, Anya Chen, and Spencer Ball. CUAD: An expert-annotated NLP dataset for legal contract review. In *NeurIPS Datasets and Benchmarks Track*, 2021. URL <https://arxiv.org/abs/2103.06268>.
- [14] Toshihide Ibaraki and Tiko Kameda. On the optimal nesting order for computing N-relational joins. *ACM Transactions on Database Systems*, 9(3):482–502, 1984. doi: 10.1145/1270.1498.
- [15] Pranab Islam, Anand Kannappan, Douwe Kiela, Rebecca Qian, Nino Scherrer, and Bertie Vidgen. FinanceBench: A new benchmark for financial question answering. *arXiv preprint arXiv:2311.11944*, 2023. URL <https://arxiv.org/abs/2311.11944>.
- [16] Carlos E. Jimenez, John Yang, Alexander Wettig, Shunyu Yao, Kexin Pei, Ofir Press, and Karthik Narasimhan. SWE-bench: Can language models resolve real-world GitHub issues? In *International Conference on Learning Representations (ICLR)*, 2024. URL <https://arxiv.org/abs/2310.06770>.
- [17] Saehan Jo and Immanuel Trummer. ThalamusDB: Approximate query processing on multi-modal data. *Proceedings of the ACM on Management of Data*, 2(3), 2024. doi: 10.1145/3654989.

- [18] Omar Khattab, Arnav Singhvi, Paridhi Maheshwari, Zhiyuan Zhang, Keshav Santhanam, Sri Vardhamanan, Saiful Haq, Ashutosh Sharma, Thomas T. Joshi, Hanna Moazam, Heather Miller, Matei Zaharia, and Christopher Potts. DSPy: Compiling declarative language model calls into self-improving pipelines. *arXiv preprint arXiv:2310.03714*, 2023. URL <https://arxiv.org/abs/2310.03714>.
- [19] Stephen C. Kleene. On notation for ordinal numbers. *Journal of Symbolic Logic*, 3(4): 150–155, 1938. doi: 10.2307/2267778.
- [20] Stephen C. Kleene. *Introduction to Metamathematics*. North-Holland, 1952.
- [21] Stephen C. Kleene. Representation of events in nerve nets and finite automata. In Claude E. Shannon and John McCarthy, editors, *Automata Studies*, pages 3–41. Princeton University Press, 1956.
- [22] Ziyang Li, Jiani Huang, and Mayur Naik. Scallop: A language for neurosymbolic programming. *Proceedings of the ACM on Programming Languages*, 7(PLDI), 2023. doi: 10.1145/3591280.
- [23] Chunwei Liu, Matthew Russo, Michael Cafarella, Lei Cao, Peter Baile Chen, Zui Chen, Michael Franklin, Tim Kraska, Samuel Madden, and Gerardo Vitagliano. A declarative system for optimizing AI workloads. In *Conference on Innovative Data Systems Research (CIDR)*, 2025. URL <https://arxiv.org/abs/2405.14696>. arXiv:2405.14696.
- [24] Shicheng Liu, Jialiang Xu, Wesley Tjangnaka, Sina J. Semnani, Chen Jie Yu, and Monica S. Lam. SUQL: Conversational search over structured and unstructured data with large language models. In *Findings of the Association for Computational Linguistics: NAACL 2024*, 2024. URL <https://arxiv.org/abs/2311.09818>.
- [25] Shu Liu, Asim Biswal, Amog Kamsetty, Audrey Cheng, Luis Mo, Lei Cao, and Ion Stoica. Optimizing LLM queries in relational data analytics workloads. *arXiv preprint arXiv:2403.05821*, 2024. URL <https://arxiv.org/abs/2403.05821>.
- [26] Mike A. Merrill, Alexander G. Shaw, Nicholas Carlini, Boxuan Li, et al. Terminal-bench: Benchmarking agents on hard, realistic tasks in command line interfaces. *arXiv preprint arXiv:2601.11868*, 2026. URL <https://arxiv.org/abs/2601.11868>.
- [27] David Mitchell, Bart Selman, and Hector Levesque. Hard and easy distributions of SAT problems. In *Proceedings of the 10th National Conference on Artificial Intelligence (AAAI)*, pages 459–465, 1992.
- [28] Guido Moerkotte and Thomas Neumann. Analysis of two existing and one new dynamic programming algorithm for the generation of optimal bushy join trees without cross products. In *Proceedings of the 32nd International Conference on Very Large Data Bases (VLDB)*, pages 930–941, 2006.
- [29] Liana Patel, Siddharth Jha, Melissa Pan, Harshit Gupta, Parth Asawa, Carlos Guestrin, and Matei Zaharia. Semantic operators: A declarative model for rich, AI-based data processing. In *Proceedings of the VLDB Endowment*, 2025. URL <https://arxiv.org/abs/2407.11418>. arXiv:2407.11418.
- [30] Mark Raasveldt and Hannes Mühleisen. DuckDB: an embeddable analytical database. In *Proceedings of the 2019 International Conference on Management of Data (SIGMOD)*, pages 1981–1984, 2019. doi: 10.1145/3299869.3320212.

- [31] Matthew Russo, Chunwei Liu, Sivaprasad Sudhir, Gerardo Vitagliano, Samuel Madden, Michael Cafarella, and Michael Franklin. Abacus: A cost-based optimizer for semantic operator systems. *arXiv preprint arXiv:2505.14661*, 2025. URL <https://arxiv.org/abs/2505.14661>.
- [32] Mohammed Saeed, Nicola De Cao, and Paolo Papotti. Querying large language models with SQL. *arXiv preprint arXiv:2304.00472*, 2023. URL <https://arxiv.org/abs/2304.00472>.
- [33] P. Griffiths Selinger, M. M. Astrahan, D. D. Chamberlin, R. A. Lorie, and T. G. Price. Access path selection in a relational database management system. In *Proceedings of the 1979 ACM SIGMOD International Conference on Management of Data*, pages 23–34, 1979. doi: 10.1145/582095.582099.
- [34] Armando Solar-Lezama, Liviu Tancau, Rastislav Bodík, Sanjit Seshia, and Vijay Saraswat. Combinatorial sketching for finite programs. In *Proceedings of the 12th International Conference on Architectural Support for Programming Languages and Operating Systems (ASPLOS)*, pages 404–415, 2006. doi: 10.1145/1168857.1168907.
- [35] Zafir Stojanovski, Oliver Stanley, Joe Sharratt, Richard Jones, et al. Reasoning gym: Reasoning environments for reinforcement learning with verifiable rewards. *arXiv preprint arXiv:2505.24760*, 2025. URL <https://arxiv.org/abs/2505.24760>.
- [36] Yunxiang Su, Tianjing Zeng, Zhongjun Ding, Yin Lin, et al. Large language model-enhanced relational operators: Taxonomy, benchmark, and analysis. *arXiv preprint arXiv:2603.02537*, 2026. URL <https://arxiv.org/abs/2603.02537>.
- [37] Matthias Urban and Carsten Binnig. CAESURA: Language models as multi-modal query planners. *arXiv preprint arXiv:2308.03424*, 2023. URL <https://arxiv.org/abs/2308.03424>.
- [38] Moshe Y. Vardi. The complexity of relational query languages. In *Proceedings of the 14th Annual ACM Symposium on Theory of Computing (STOC)*, pages 137–146, 1982. doi: 10.1145/800070.802186.
- [39] Daren Wang. Think, but don’t overthink: Reproducing recursive language models. *arXiv preprint arXiv:2603.02615*, 2026. URL <https://arxiv.org/abs/2603.02615>.
- [40] Xingyao Wang, Yangyi Chen, Lifan Yuan, Yizhe Zhang, Yunzhu Peng, and Heng Ji. Executable code actions elicit better LLM agents. In *International Conference on Machine Learning (ICML)*, 2024. URL <https://arxiv.org/abs/2402.01030>.
- [41] Shunyu Yao, Jeffrey Zhao, Dian Yu, Nan Du, Izhak Shafran, Karthik Narasimhan, and Yuan Cao. ReAct: Synergizing reasoning and acting in language models. In *International Conference on Learning Representations (ICLR)*, 2023. URL <https://arxiv.org/abs/2210.03629>.
- [42] Alex L. Zhang, Tim Kraska, and Omar Khattab. Recursive language models. *arXiv preprint arXiv:2512.24601*, 2025. URL <https://arxiv.org/abs/2512.24601>.
- [43] Andrew Zhao, Yiran Wu, Yang Yue, Tong Wu, Quentin Xu, Yang Yue, Matthieu Lin, Shenzi Wang, Qingyun Wu, Zilong Zheng, and Gao Huang. Absolute zero: Reinforced self-play reasoning with zero data. *arXiv preprint arXiv:2505.03335*, 2025. URL <https://arxiv.org/abs/2505.03335>.